

```
In [169]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [170]: # Load the dataset
Telco_df = pd.read_csv('Telco-Customer-Churn.csv')
```

```
In [171]: print(Telco_df.head()) # Display the first few rows
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
0	7590-VHVEG	Female	0	Yes	No	1	No
1	5575-GNVDE	Male	0	No	No	34	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes
3	7795-CFOCW	Male	0	No	No	45	No
4	9237-HQITU	Female	0	No	No	2	Yes

	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection
0	No phone service	DSL	No	...	No
1	No	DSL	Yes	...	Yes
2	No	DSL	Yes	...	No
3	No phone service	DSL	Yes	...	Yes
4	No	Fiber optic	No	...	No

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling
0	No	No	No	Month-to-month	Yes
1	No	No	No	One year	No
2	No	No	No	Month-to-month	Yes
3	Yes	No	No	One year	No
4	No	No	No	Month-to-month	Yes

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

```
In [172]: print (Telco_df.shape)
          Telco_df.isnull().sum()
```

```
(7043, 21)
```

```
Out[172]: customerID      0
          gender          0
          SeniorCitizen    0
          Partner          0
          Dependents       0
          tenure          0
          PhoneService     0
          MultipleLines    0
          InternetService  0
          OnlineSecurity   0
          OnlineBackup     0
          DeviceProtection 0
          TechSupport      0
          StreamingTV      0
          StreamingMovies  0
          Contract        0
          PaperlessBilling 0
          PaymentMethod    0
          MonthlyCharges   0
          TotalCharges     0
          Churn            0
          dtype: int64
```

```
In [173]: print(Telco_df.info()) # Summary of the dataset
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
None
```

```
In [174]: # Check for duplicate rows
Telco_df.duplicated().sum()
```


```
Out[174]: 0
```

VISUALIZATION


```
In [175]: # Visualization 1: Churn distribution
sns.countplot(x='Churn', data=Telco_df)
plt.title('Churn Distribution')
plt.show()
```



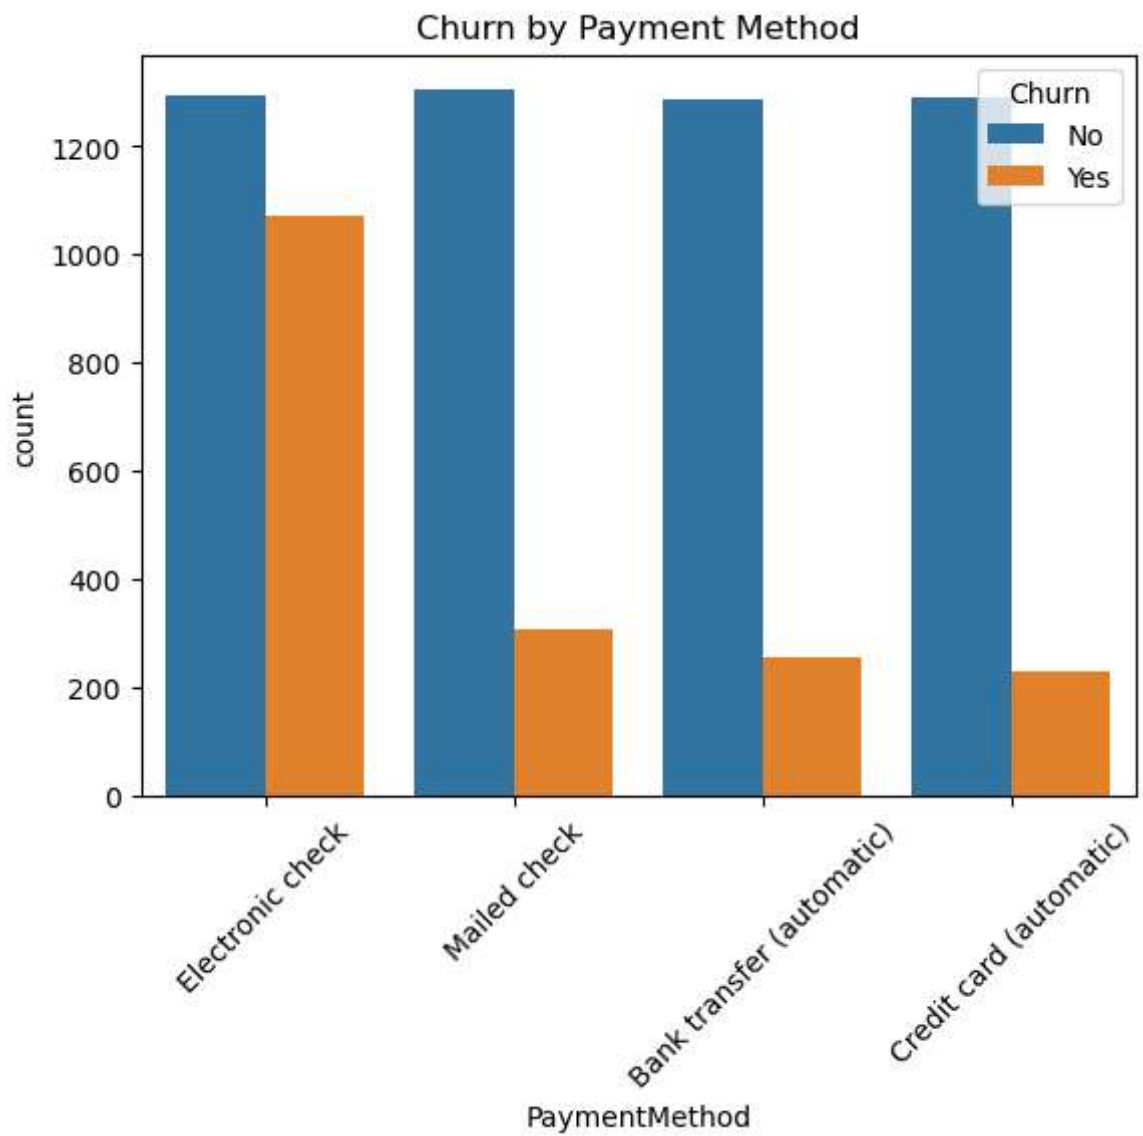
```
In [176]: # Visualization 2: Churn by Contract type
sns.countplot(x='Contract', hue='Churn', data=Telco_df)
plt.title('Churn by Contract Type')
plt.show()
```




```
In [177]: # Visualization 3: Churn by Internet Service type
sns.countplot(x='InternetService', hue='Churn', data=Telco_df)
plt.title('Churn by Internet Service Type')
plt.show()
```



```
In [178]: # Visualization 4: Churn by Payment Method
sns.countplot(x='PaymentMethod', hue='Churn', data=Telco_df)
plt.title('Churn by Payment Method')
plt.xticks(rotation=45)
plt.show()
```




```
In [179]: # Visualization 5: Churn by Tenure
sns.boxplot(x='Churn', y='tenure', data=Telco_df)
plt.title('Churn by Tenure')
plt.show()
```




so we are not taking any action to handle the outliers,because a customer can be in the company for many months. Even though the customer had been in the company for a long time, we should consider them as part of our analysis.


```
In [180]: # 6. Pie chart for churn distribution
plt.figure(figsize=(6, 6))
plt.pie(Telco_df['Churn'].value_counts(), labels=['No', 'Yes'], autopct='%1.1f%%')
plt.title('Churn Distribution')
plt.show()
```



```
In [181]: # Visualization 7: Correlation Heatmap
correlation_matrix = Telco_df.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```



DATA CLEANING

```
In [182]: # Assuming 'No internet service' means the customer does not have that service
cols_fillna = ['MultipleLines', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
               'StreamingTV', 'StreamingMovies']
Telco_df[cols_fillna] = Telco_df[cols_fillna].replace('No internet service', '')
```

```
In [183]: # Assuming 'No internet service' means the customer does not have that service
cols_fillna = ['MultipleLines', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
               'StreamingTV', 'StreamingMovies']
Telco_df[cols_fillna] = Telco_df[cols_fillna].replace('No phone service', 'No')
```

```
In [184]: Telco_df
```

Out[184]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLir
0	7590-VHVEG	Female	0	Yes	No	1	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	
4	9237-HQITU	Female	0	No	No	2	Yes	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	\
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	\
7040	4801-JZAZL	Female	0	Yes	Yes	11	No	
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	\
7042	3186-AJIEK	Male	0	No	No	66	Yes	

7043 rows × 21 columns



```
In [185]: # Drop irrelevant columns like customerID as it does not contribute to the pre
Telco_df.drop(columns=['customerID'], inplace=True)
```

Feature Engineering

```
In [186]: # Convert 'Churn' column to binary values
Telco_df['Churn'] = Telco_df['Churn'].map({'Yes': 1, 'No': 0})
```

```
In [187]: # Create binary features for 'Partner', 'Dependents', 'PhoneService', 'PaperlessBilling'
Telco_df['Partner'] = Telco_df['Partner'].map({'Yes': 1, 'No': 0})
Telco_df['Dependents'] = Telco_df['Dependents'].map({'Yes': 1, 'No': 0})
Telco_df['PhoneService'] = Telco_df['PhoneService'].map({'Yes': 1, 'No': 0})
Telco_df['PaperlessBilling'] = Telco_df['PaperlessBilling'].map({'Yes': 1, 'No': 0})
```

```
In [188]: Telco_df
```

Out[188]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	0	1	0	1	0	No	
1	Male	0	0	0	34	1	No	
2	Male	0	0	0	2	1	No	
3	Male	0	0	0	45	0	No	
4	Female	0	0	0	2	1	No	Fiber
...	
7038	Male	0	1	1	24	1	Yes	
7039	Female	0	1	1	72	1	Yes	Fiber
7040	Female	0	1	1	11	0	No	
7041	Male	1	1	0	4	1	Yes	Fiber
7042	Male	0	0	0	66	1	No	Fiber

7043 rows × 20 columns

```
Telco_df
```

Out[189]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
0	Female	0	1	0	1	0	No	
1	Male	0	0	0	34	1	No	
2	Male	0	0	0	2	1	No	
3	Male	0	0	0	45	0	No	
4	Female	0	0	0	2	1	No	FiberOptic
...	
7038	Male	0	1	1	24	1	Yes	
7039	Female	0	1	1	72	1	Yes	FiberOptic
7040	Female	0	1	1	11	0	No	
7041	Male	1	1	0	4	1	Yes	FiberOptic
7042	Male	0	0	0	66	1	No	FiberOptic

7043 rows × 20 columns



```
In [190]: Telco_df = pd.get_dummies(Telco_df, columns=['gender', 'MultipleLines', 'Inter',
'OnlineSecurity', 'OnlineBackup',
'TechSupport', 'StreamingTV', 'Str
'Contract'], drop_first=True)
```

In [191]: Telco_df

Out[191]:

	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling	PaymentMetho
0	0	1	0	1	0	1	Electronic chec
1	0	0	0	34	1	0	Mailed chec
2	0	0	0	2	1	1	Mailed chec
3	0	0	0	45	0	0	Bank transfe (automatic
4	0	0	0	2	1	1	Electronic chec
...
7038	0	1	1	24	1	1	Mailed chec
7039	0	1	1	72	1	1	Credit car (automatic
7040	0	1	1	11	0	1	Electronic chec
7041	1	1	0	4	1	1	Mailed chec
7042	0	0	0	66	1	1	Bank transfe (automatic

7043 rows × 22 columns



```
In [192]: # Convert TotalCharges to numeric
Telco_df['TotalCharges'] = pd.to_numeric(Telco_df['TotalCharges'], errors='coe
```

```
In [193]: #calculate the ratio of MonthlyCharges to TotalCharges to see the average mont
Telco_df['AvgMonthlySpending'] = Telco_df['TotalCharges'] / Telco_df['tenure']
```

In [194]: Telco_df

Out[194]:

	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling	PaymentMetho
0	0	1	0	1	0	1	Electronic chec
1	0	0	0	34	1	0	Mailed chec
2	0	0	0	2	1	1	Mailed chec
3	0	0	0	45	0	0	Bank transfe (automatic
4	0	0	0	2	1	1	Electronic chec
...
7038	0	1	1	24	1	1	Mailed chec
7039	0	1	1	72	1	1	Credit car (automatic
7040	0	1	1	11	0	1	Electronic chec
7041	1	1	0	4	1	1	Mailed chec
7042	0	0	0	66	1	1	Bank transfe (automatic

7043 rows × 23 columns



Feature Scaling

```
In [195]: # Feature Scaling (if required)
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
numerical_features = ['tenure', 'MonthlyCharges', 'TotalCharges']
Telco_df[numerical_features] = scaler.fit_transform(Telco_df[numerical_features])
```

In [196]: Telco_df

Out[196]:

	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling	PaymentMethod
0	0	1	0	-1.277445	0	1	Electronic c
1	0	0	0	0.066327	1	0	Mailed c
2	0	0	0	-1.236724	1	1	Mailed c
3	0	0	0	0.514251	0	0	Bank tra (autorr
4	0	0	0	-1.236724	1	1	Electronic c
...
7038	0	1	1	-0.340876	1	1	Mailed c
7039	0	1	1	1.613701	1	1	Credit (autorr
7040	0	1	1	-0.870241	0	1	Electronic c
7041	1	1	0	-1.155283	1	1	Mailed c
7042	0	0	0	1.369379	1	1	Bank tra (autorr

7043 rows × 23 columns



In [197]: *# Check for missing values*
Telco_df.isnull().sum()

Out[197]: SeniorCitizen 0
Partner 0
Dependents 0
tenure 0
PhoneService 0
PaperlessBilling 0
PaymentMethod 0
MonthlyCharges 0
TotalCharges 11
Churn 0
gender_Male 0
MultipleLines_Yes 0
InternetService_Fiber optic 0
InternetService_No 0
OnlineSecurity_Yes 0
OnlineBackup_Yes 0
DeviceProtection_Yes 0
TechSupport_Yes 0
StreamingTV_Yes 0
StreamingMovies_Yes 0
Contract_One year 0
Contract_Two year 0
AvgMonthlySpending 11
dtype: int64

DEALING WITH MISSING VALUES

```
In [198]: total_charges_mean = Telco_df['TotalCharges'].mean()  
#Replace the missing values in the 'TotalCharges' column with the calculated mean  
Telco_df['TotalCharges'].fillna(total_charges_mean, inplace=True)
```

```
In [199]: #dealing with the missing values using the mean  
avg_monthly_spending_mean = Telco_df['AvgMonthlySpending'].mean()  
#Replace the missing values in the 'AvgMonthlySpending' column with the calculated mean  
Telco_df['AvgMonthlySpending'].fillna(avg_monthly_spending_mean, inplace=True)
```

```
In [200]: Telco_df.isnull().sum()
```

```
Out[200]: SeniorCitizen      0  
Partner      0  
Dependents    0  
tenure        0  
PhoneService  0  
PaperlessBilling  0  
PaymentMethod  0  
MonthlyCharges  0  
TotalCharges  0  
Churn          0  
gender_Male    0  
MultipleLines_Yes  0  
InternetService_Fiber optic  0  
InternetService_No  0  
OnlineSecurity_Yes  0  
OnlineBackup_Yes  0  
DeviceProtection_Yes  0  
TechSupport_Yes  0  
StreamingTV_Yes  0  
StreamingMovies_Yes  0  
Contract_One year  0  
Contract_Two year  0  
AvgMonthlySpending  0  
dtype: int64
```

VALIDATION SPLIT

```
In [201]: from sklearn.model_selection import train_test_split  
from sklearn.neural_network import MLPClassifier  
from sklearn.preprocessing import LabelEncoder  
from sklearn.metrics import classification_report, confusion_matrix  
from dmbs import classificationSummary
```

```
In [202]: data into features (X) and target (y)
'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'PaperlessB
Churn']
```

```
In [203]: # Splitting the dataset into training and testing sets
train_X, valid_X, train_y, valid_y = train_test_split(X, y, test_size=0.2, ran
```

NEURAL NETWORK

```
In [204]: clf = MLPClassifier(hidden_layer_sizes=(200, 100), activation='logistic', solv
```

```
In [205]: clf.fit(train_X, train_y.values)
```

```
Out[205]: MLPClassifier(activation='logistic', batch_size=256,
                        hidden_layer_sizes=(200, 100), max_iter=2000)
```

```
In [206]: clf.predict(X)
```

```
Out[206]: array([0, 0, 0, ..., 0, 1, 0], dtype=int64)
```

```
In [207]: #Network structure  
print('Intercepts')  
print(clf.intercepts_)
```

Intercepts

```
[array([-0.07711076,  0.04778391, -0.11437006,  0.16322386, -0.05399742,
        0.06280103, -0.06295765, -0.13327462, -0.07654465, -0.0741688 ,
        0.05018679, -0.01543338, -0.12097818, -0.08782992, -0.06460622,
        -0.08328986, -0.02982481, -0.01698305,  0.13822986, -0.0878224 ,
        0.07614872,  0.10424416,  0.02954134,  0.04933228,  0.14071913,
        0.04604924, -0.12765704,  0.09622789,  0.04664133,  0.04145939,
        0.06980371,  0.07832422,  0.10275634,  0.01520635, -0.12026865,
        -0.06718483, -0.05889039, -0.08507691, -0.12489526,  0.06690968,
        -0.04913522,  0.11136058, -0.06579893,  0.05574222,  0.07085085,
        -0.09586987, -0.07378966, -0.09442997,  0.01864218, -0.04196119,
        0.08233836,  0.09994046,  0.10435348, -0.01547792,  0.08910943,
        -0.14863354, -0.04997154, -0.05405346,  0.10181535,  0.07388298,
        0.07375984,  0.09568427, -0.02953724, -0.01568243, -0.00917402,
        0.06416815,  0.06843717,  0.03059337, -0.08253441,  0.05186856,
        0.05507195,  0.03297723, -0.09402142,  0.0479809 ,  0.02748668,
        0.01661871,  0.01167706,  0.04206024,  0.06373487,  0.14502593,
        -0.04502546,  0.16194255,  0.06799084,  0.01056557,  0.09549596,
        -0.05057188, -0.07249029, -0.06127656,  0.08602927, -0.03266914,
        0.03470659, -0.09969943,  0.10521663, -0.04670525, -0.0260195 ,
        -0.00987203,  0.03982132,  0.06858372,  0.07712721,  0.06821213,
        -0.05507857, -0.03856469,  0.0637783 , -0.10458279, -0.00100946,
        -0.12456497,  0.11352142,  0.01896289,  0.14389551,  0.07789877,
        0.00304002, -0.10607939,  0.03724744, -0.09502081, -0.0121915 ,
        0.01661286, -0.08669511, -0.05088586,  0.03881723,  0.0180952 ,
        0.1437362 , -0.09645999, -0.01274096, -0.00719962,  0.03376839,
        0.02546789, -0.12776966,  0.12671569,  0.05497335, -0.05079617,
        -0.0181749 ,  0.11890329,  0.04643214,  0.0549724 ,  0.0188724 ,
        -0.057389 , -0.02018207,  0.00730316,  0.11797376,  0.05259888,
        -0.0557688 , -0.08715235, -0.07190564, -0.05273055,  0.01922235,
        -0.00453224,  0.11673014,  0.08466116,  0.1150289 ,  0.04381728,
        -0.01404847, -0.07236474, -0.08352699,  0.11910235, -0.08244437,
        0.05187741,  0.03364077,  0.09214004,  0.03660464,  0.04965636,
        -0.06029192,  0.08509258, -0.05481139, -0.06163182, -0.10427193,
        0.02788885,  0.04705745,  0.05053703,  0.0755811 ,  0.09546765,
        0.00078261, -0.08677788, -0.05073877,  0.04878723,  0.11729899,
        0.0511977 ,  0.03651654, -0.0076614 ,  0.06618771,  0.12390693,
        0.03620347, -0.01704678,  0.10710408,  0.07250365, -0.03203895,
        -0.1169381 ,  0.05736459, -0.0535965 ,  0.02953488,  0.03588595,
        0.02550196, -0.09880053, -0.01919622, -0.03628041, -0.08358034,
        0.04978479,  0.13139801, -0.07146026,  0.02972915, -0.00810207]), arr
ay([ 5.32295116e-02, -8.32804305e-02,  2.10037985e-03,  4.53619117e-02,
    -9.14240261e-02, -7.31747962e-02,  6.10768135e-02, -8.30551147e-02,
    -3.70508351e-02,  5.03482994e-02,  1.13100221e-02, -7.94008044e-02,
    -7.10224817e-02, -7.21744390e-03, -4.76103620e-02,  4.00430714e-02,
    -4.25554549e-02,  5.78202683e-02, -6.28524628e-02, -5.38364907e-02,
    -6.92670029e-02, -1.12842435e-01, -1.39195298e-02, -5.78254708e-02,
     3.63953678e-02, -9.04994729e-02, -1.00320687e-02, -2.07142544e-02,
    -4.42791536e-02, -7.56208117e-02,  7.68847492e-02, -7.29476919e-02,
    -1.05801169e-01, -4.95628185e-02,  4.87847088e-02, -9.71934795e-02,
     3.88843209e-02, -1.08549199e-02, -2.20046808e-03,  4.27585415e-02,
     4.22117290e-02, -7.69699988e-02, -6.14362826e-02,  2.05567835e-02,
    -1.56535577e-02, -9.41029854e-03, -1.86458573e-02,  5.08265956e-02,
     5.43924321e-02,  2.51265149e-02, -6.77285137e-02, -8.55760263e-03,
    -7.45661044e-02,  3.83314009e-02,  1.75649614e-02,  1.60116824e-03,
     7.64290969e-02,  1.20938913e-03, -4.81318841e-02, -4.49948585e-02,
     5.42117901e-02, -5.65470492e-02, -3.74689535e-03, -5.52694032e-02,
```

```

2.27156884e-02, -1.04961403e-01, -3.23709447e-02, -4.92558876e-02,
2.13948375e-03, -7.97998621e-02, -2.55598297e-02, 8.38673717e-02,
-1.69493803e-03, -2.39888721e-02, -9.73972094e-02, -2.92642318e-03,
7.93858524e-03, -5.23213698e-02, 3.34637271e-02, -2.42982546e-02,
5.70511069e-02, -4.16644537e-02, -6.34847691e-03, -7.02037905e-02,
2.45368582e-02, -1.03574353e-01, -3.19790869e-02, -4.10134984e-02,
7.81284562e-02, -4.51689461e-02, -3.72090685e-02, -6.41228594e-02,
-8.50386348e-02, -5.33562435e-02, 8.88003827e-03, -8.31653411e-02,
6.06990400e-05, 4.54496327e-02, 6.58351156e-02, -4.34009324e-02]),
array([-0.03900596])]

```

```

In [208]: # training performance
classificationSummary(train_y, clf.predict(train_X))

```

Confusion Matrix (Accuracy 0.7989)

	Prediction	
Actual	0	1
0	3850	288
1	845	651

```

In [209]: # validation performance
classificationSummary(valid_y, clf.predict(valid_X))

```

Confusion Matrix (Accuracy 0.8098)

	Prediction	
Actual	0	1
0	960	76
1	192	181

RANDOM FOREST

```

In [210]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_

```

```

In [211]: rf = RandomForestClassifier(random_state=1)
rf.fit(train_X, train_y)

```

Out[211]: RandomForestClassifier(random_state=1)

```

In [212]: rf = RandomForestClassifier(random_state=1)
rf.fit(train_X, train_y)

```

Out[212]: RandomForestClassifier(random_state=1)

```
In [213]: train_pred = rf.predict(train_X)
         valid_pred = rf.predict(valid_X)
```

```
In [214]: # Calculate Metrics for Training Data
         train_cm = confusion_matrix(train_y, train_pred)
         train_accuracy = accuracy_score(train_y, train_pred)
         train_precision = precision_score(train_y, train_pred)
         train_recall = recall_score(train_y, train_pred)
         train_f1_score = f1_score(train_y, train_pred)
```

```
In [215]: # Calculate Metrics for Validation Data
         valid_cm = confusion_matrix(valid_y, valid_pred)
         valid_accuracy = accuracy_score(valid_y, valid_pred)
         valid_precision = precision_score(valid_y, valid_pred)
         valid_recall = recall_score(valid_y, valid_pred)
         valid_f1_score = f1_score(valid_y, valid_pred)
```

```
In [216]: # Print the Metrics
         print("Random Forest Metrics:")
         print("Training Accuracy:", train_accuracy)
         print("Training Precision:", train_precision)
         print("Training Recall:", train_recall)
         print("Training F1 Score:", train_f1_score)
         print("\nValidation Accuracy:", valid_accuracy)
         print("Validation Precision:", valid_precision)
         print("Validation Recall:", valid_recall)
         print("Validation F1 Score:", valid_f1_score)
```

```
Random Forest Metrics:
Training Accuracy: 0.9978700745473909
Training Precision: 0.9986559139784946
Training Recall: 0.9933155080213903
Training F1 Score: 0.9959785522788204
```

```
Validation Accuracy: 0.7977288857345636
Validation Precision: 0.6617647058823529
Validation Recall: 0.48257372654155495
Validation F1 Score: 0.5581395348837209
```

```
In [217]: # Print the Confusion Matrices
print("\nTraining Confusion Matrix:")
print(train_cm)
print("\nValidation Confusion Matrix:")
print(valid_cm)
```

Training Confusion Matrix:

```
[[4136    2]
 [  10 1486]]
```

Validation Confusion Matrix:

```
[[944  92]
 [193 180]]
```

LOGISTIC REGRESSION

```
In [218]: #Import necessary models and evaluation metrics
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
```

```
In [219]: logreg = LogisticRegression(max_iter=1000)
logreg.fit(train_X, train_y)
y_pred_logreg = logreg.predict(valid_X)
```

```
In [220]: # Evaluate the models
print("Logistic Regression Metrics:")
print("Accuracy:", accuracy_score(valid_y, y_pred_logreg))
print("Precision:", precision_score(valid_y, y_pred_logreg))
print("Recall:", recall_score(valid_y, y_pred_logreg))
print("F1 Score:", f1_score(valid_y, y_pred_logreg))
```

Logistic Regression Metrics:

Accuracy: 0.815471965933286

Precision: 0.6805111821086262

Recall: 0.5710455764075067

F1 Score: 0.6209912536443148

In summary, based on the provided metrics, the Logistic Regression model generally performs better than the other algorithms.

```
In [ ]:
```

